




MACHINE LEARNING FOR FINANCE IN PYTHON

# Scaling data and KNN Regression

Nathan George

Data Science Professor





# Feature selection: remove weekdays

```
print(feature_names)

['10d_close_pct',
 '14-day SMA',
 '14-day RSI',
 '200-day SMA',
 '200-day RSI',
 'Adj_Volume_1d_change',
 'Adj_Volume_1d_change_SMA',
 'weekday_1',
 'weekday_2',
 'weekday_3',
 'weekday_4']


print(feature_names[:-4])


['10d_close_pct',
 '14-day SMA',
 '14-day RSI',
 '200-day SMA',
 '200-day RSI',
 'Adj_Volume_1d_change',
 'Adj_Volume_1d_change_SMA']
```




# Remove weekdays

```
train_features = train_features.iloc[:, :-4]  
test_features = test_features.iloc[:, :-4]
```











$$D(A, B) = \sum_i (|a_i - b_i|)^{(1/p)}$$




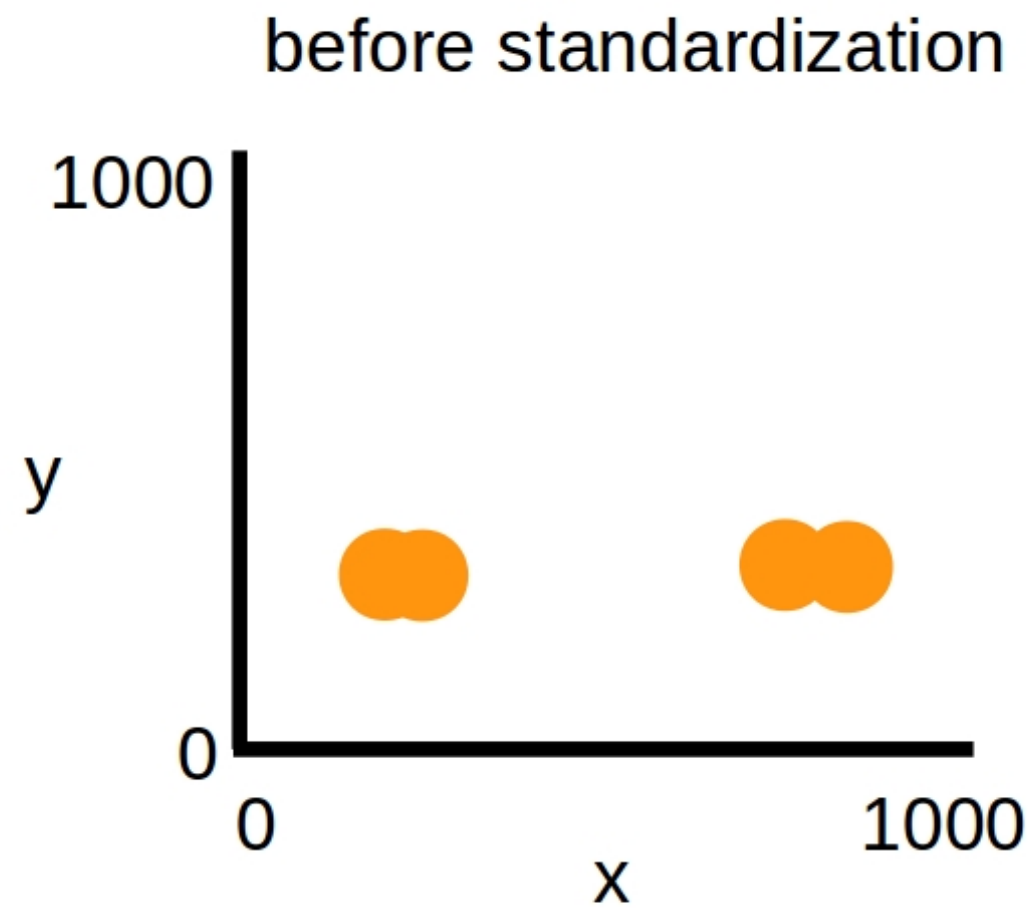




# Scaling options

Scaling options:

- min-max
- standardization
- median-MAD
- map to arbitrary function (e.g. sigmoid, tanh)







# sklearn's scaler

```
from sklearn.preprocessing import scaler

sc = scaler()
scaled_train_features = sc.fit_transform(train_features)
scaled_test_features = sc.transform(test_features)
```





# Making subplots

```
# create figure and list containing axes
f, ax = plt.subplots(nrows=2, ncols=1)

# plot histograms of before and after scaling
train_features.iloc[:, 2].hist(ax=ax[0])
ax[1].hist(scaled_train_features[:, 2])
plt.show()
```



MACHINE LEARNING FOR FINANCE IN PYTHON

**Scale data and use  
KNN!**




MACHINE LEARNING FOR FINANCE IN PYTHON

# Neural Networks

Nathan George

Data Science Professor



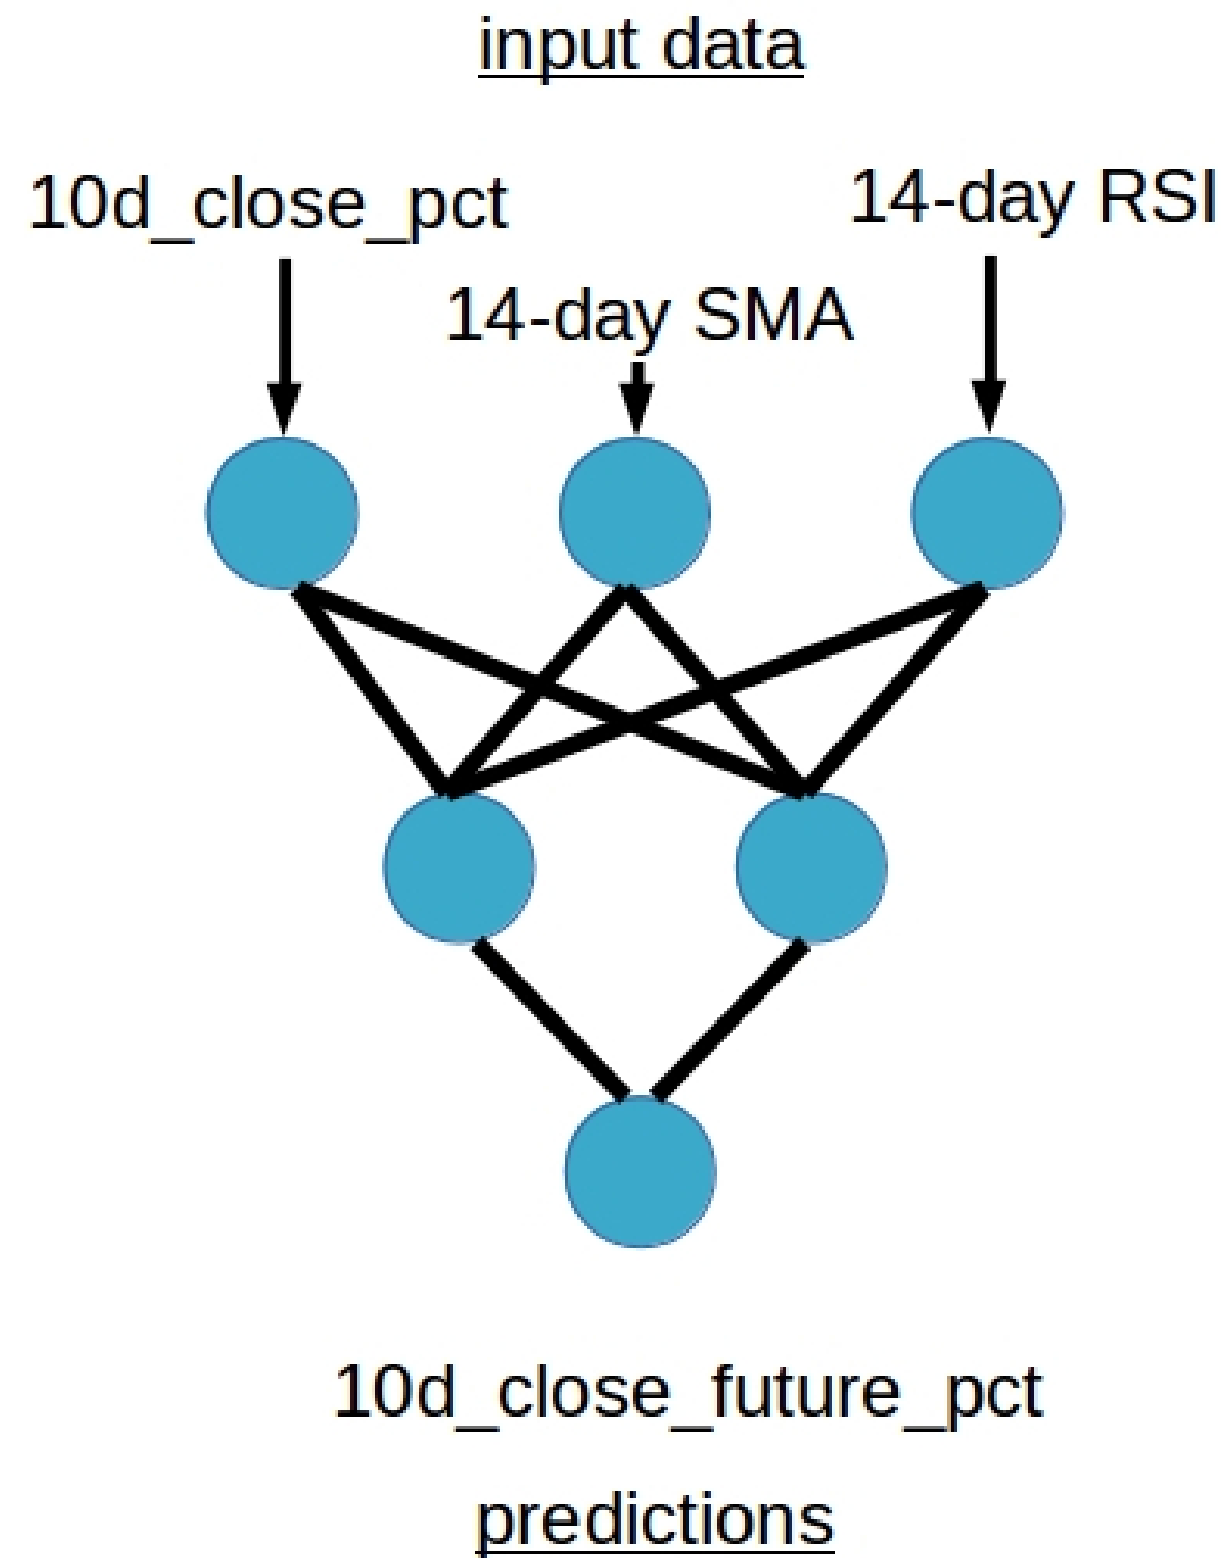


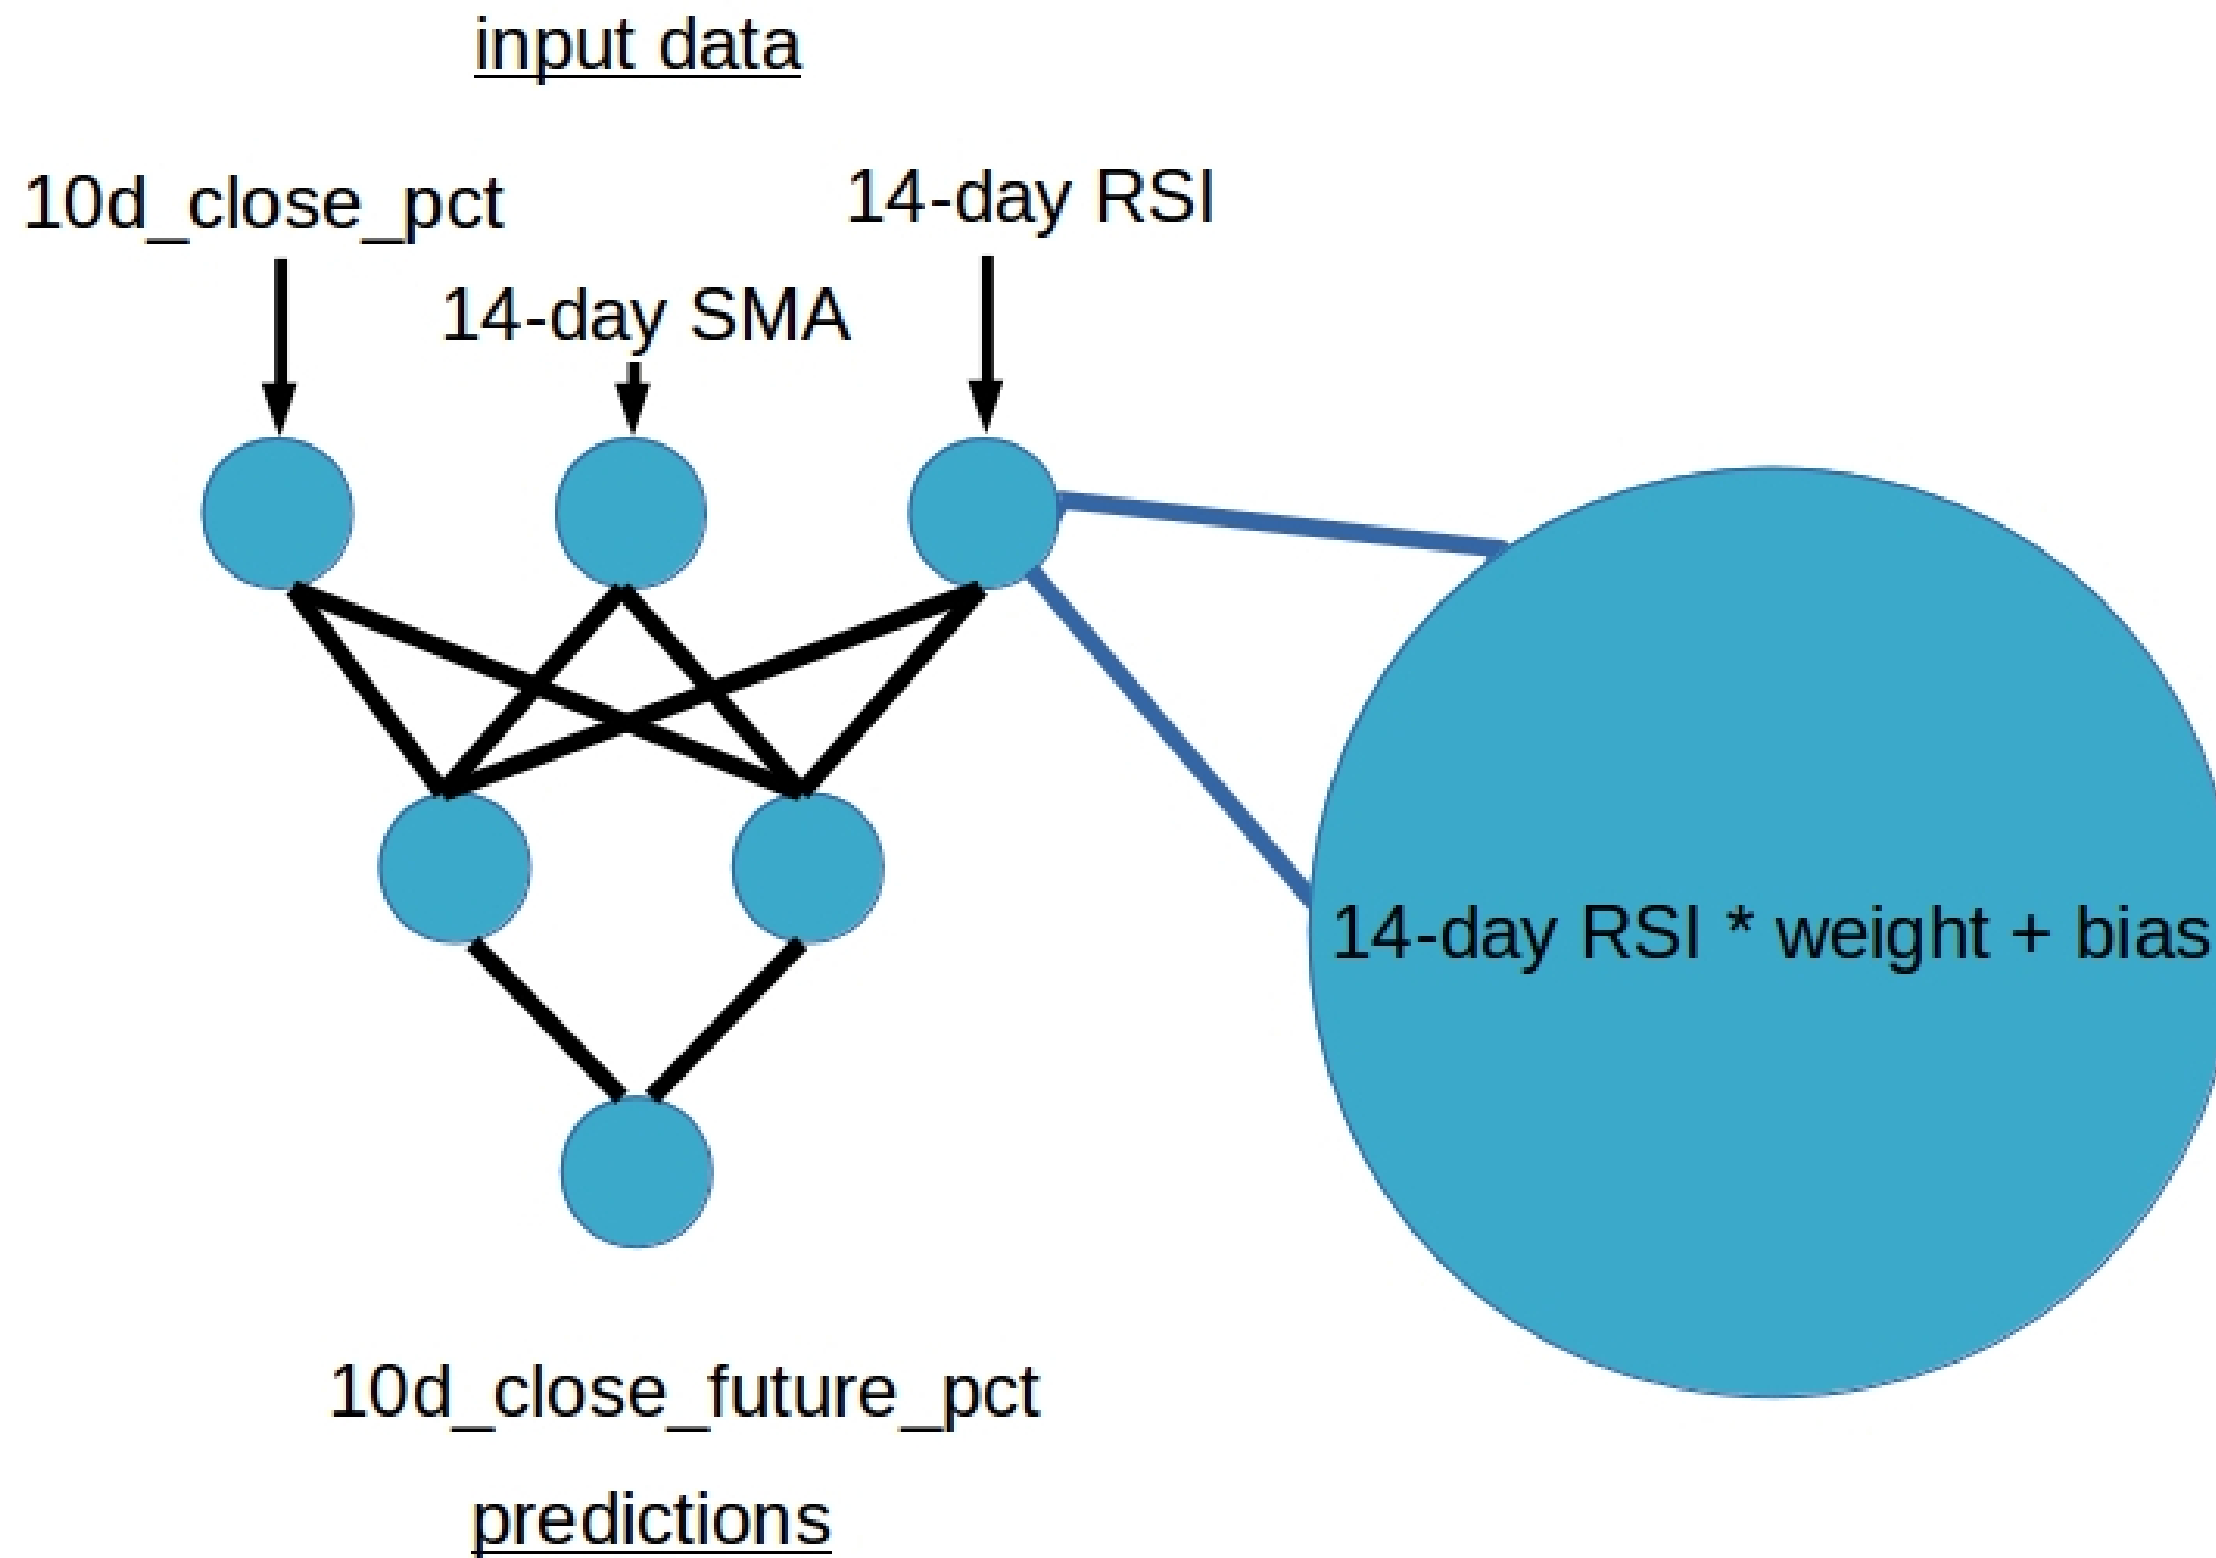


# Neural networks have potential

Neural nets have:

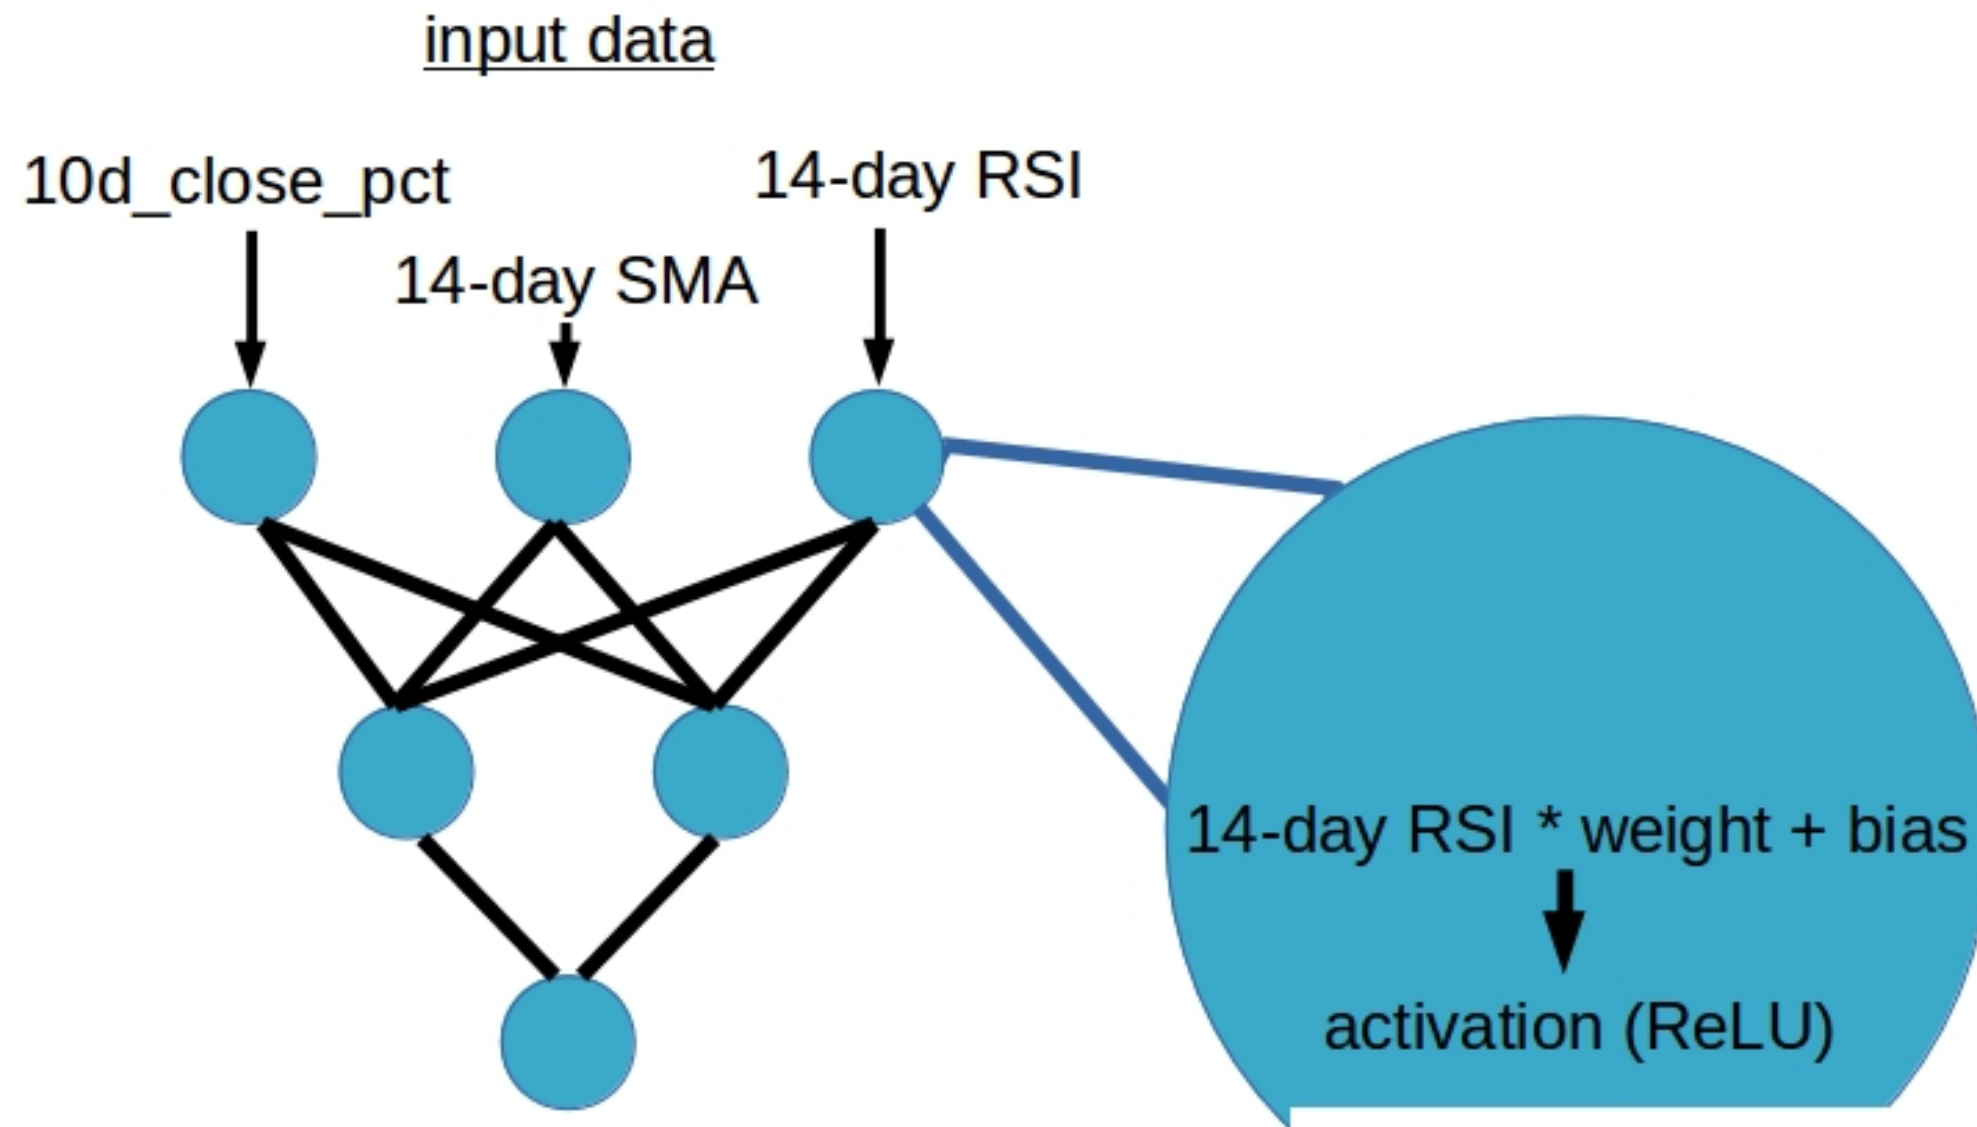
- non-linearity
- variable interactions
- customizability









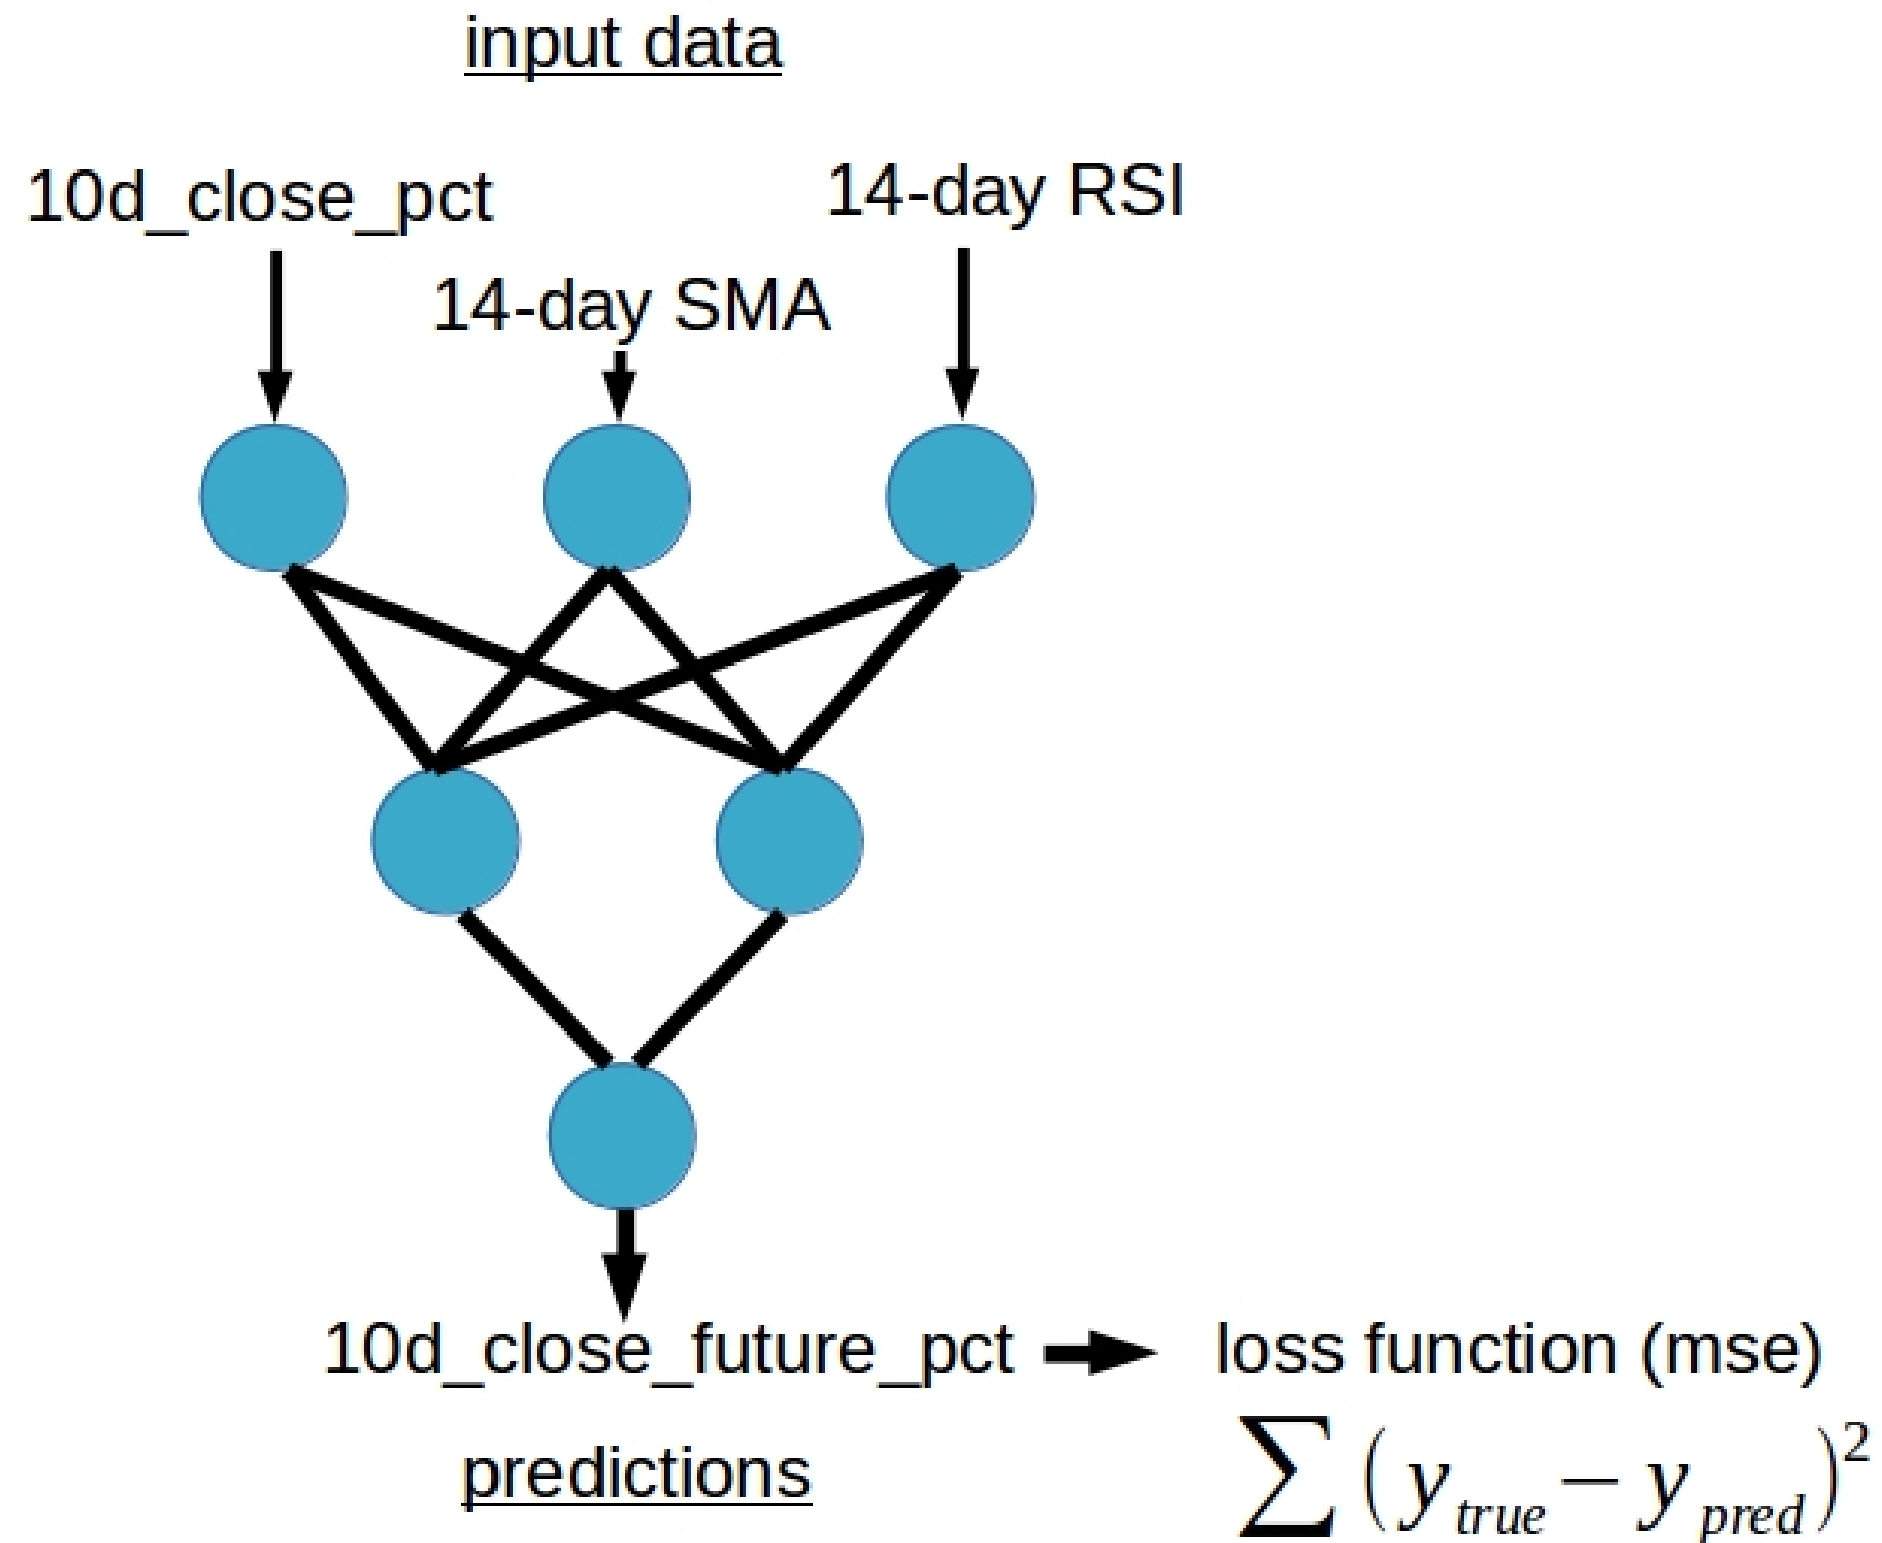
$$\sum_i w_i x_i + b$$



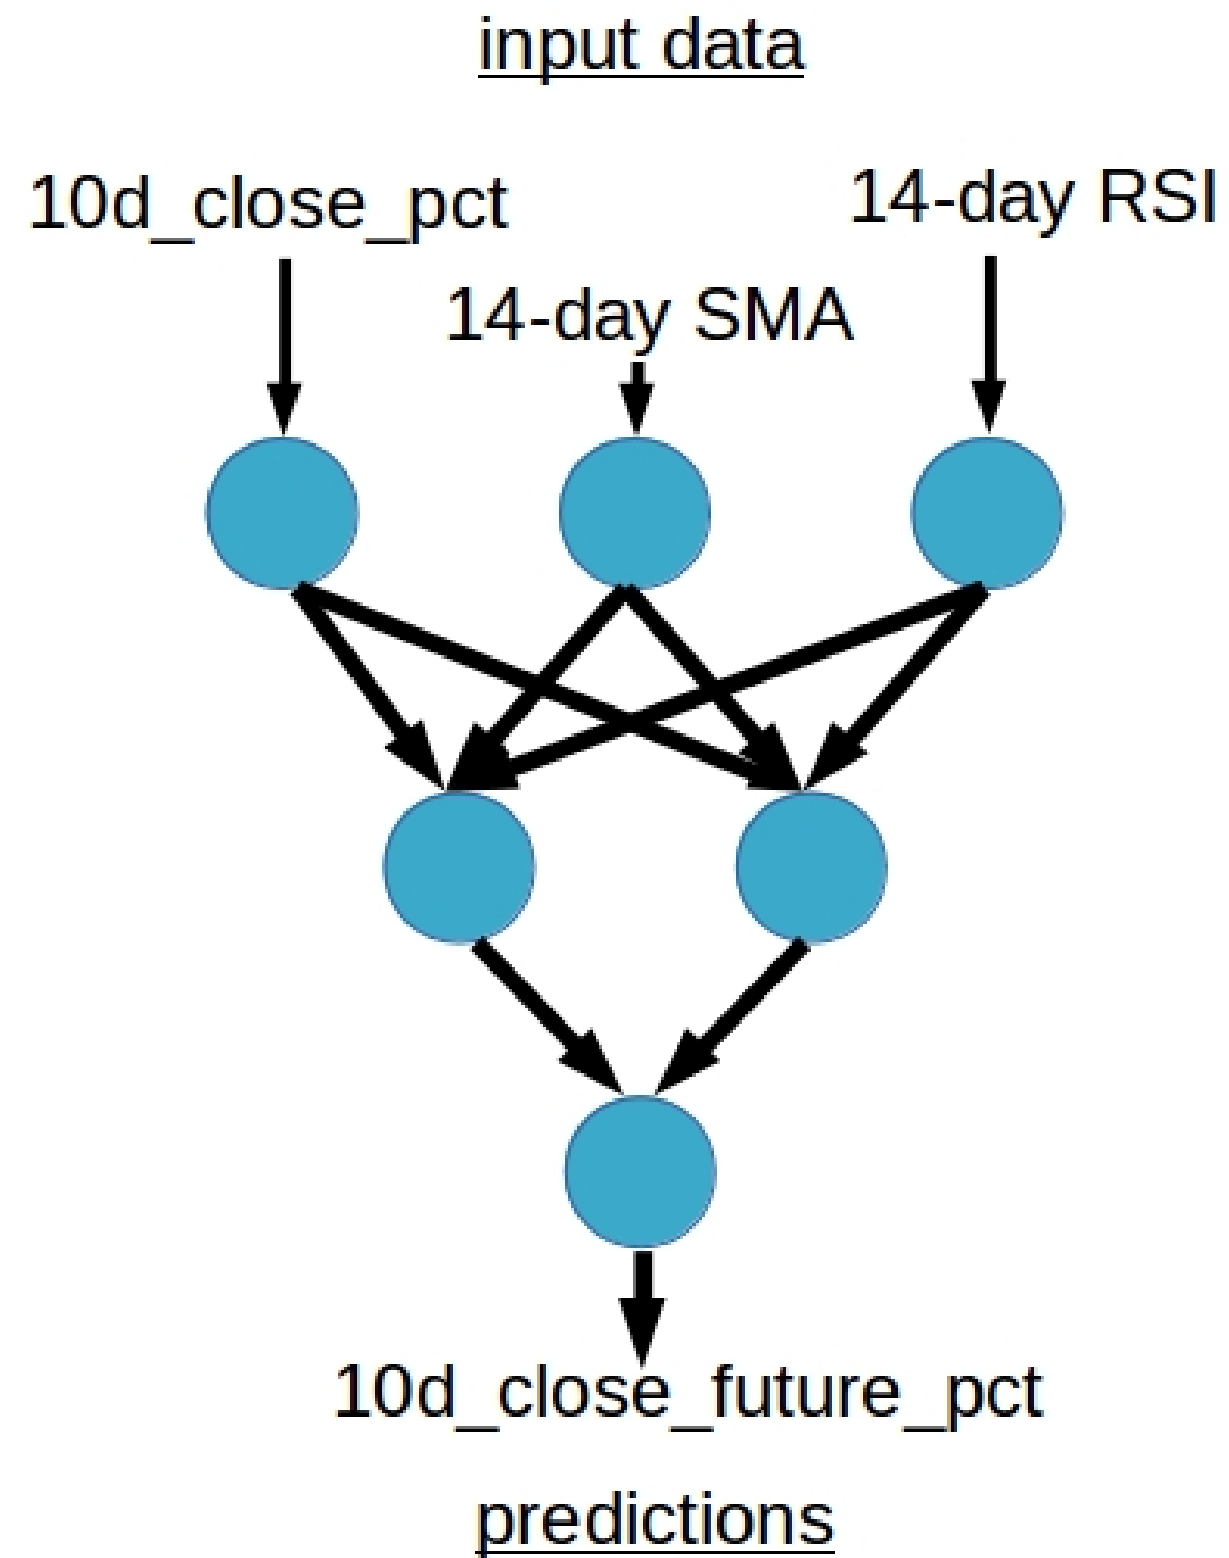
10d\_close\_future\_pct  
predictions

















# Keras



# TensorFlow



# Implementing a neural net with keras

```
from keras.models import Sequential  
from keras.layers import Dense
```



# Implementing a neural net with keras

```
from keras.models import Sequential
from keras.layers import Dense

model = Sequential()

model.add(Dense(50,
                input_dim=scaled_train_features.shape[1],
                activation='relu'))
model.add(Dense(10, activation='relu'))
model.add(Dense(1, activation='linear'))
```




# Fitting the model

```
model.compile(optimizer='adam', loss='mse')  
history = model.fit(scaled_train_features, train_targets, epochs=50)
```



# Examining the loss

```
plt.plot(history.history['loss'])  
plt.title('loss:' + str(round(history.history['loss'][-1], 6)))  
plt.xlabel('epoch')  
plt.ylabel('loss')  
plt.show()
```





# Checking out performance

```
from sklearn.metrics import r2_score

# calculate R^2 score
train_preds = model.predict(scaled_train_features)
print(r2_score(train_targets, train_preds))
```


```
0.4771387560719418
```





# Plot performance

```
# plot predictions vs actual  
plt.scatter(train_preds, train_targets)  
plt.xlabel('predictions')  
plt.ylabel('actual')  
plt.show()
```





MACHINE LEARNING FOR FINANCE IN PYTHON

**Make a neural net!**

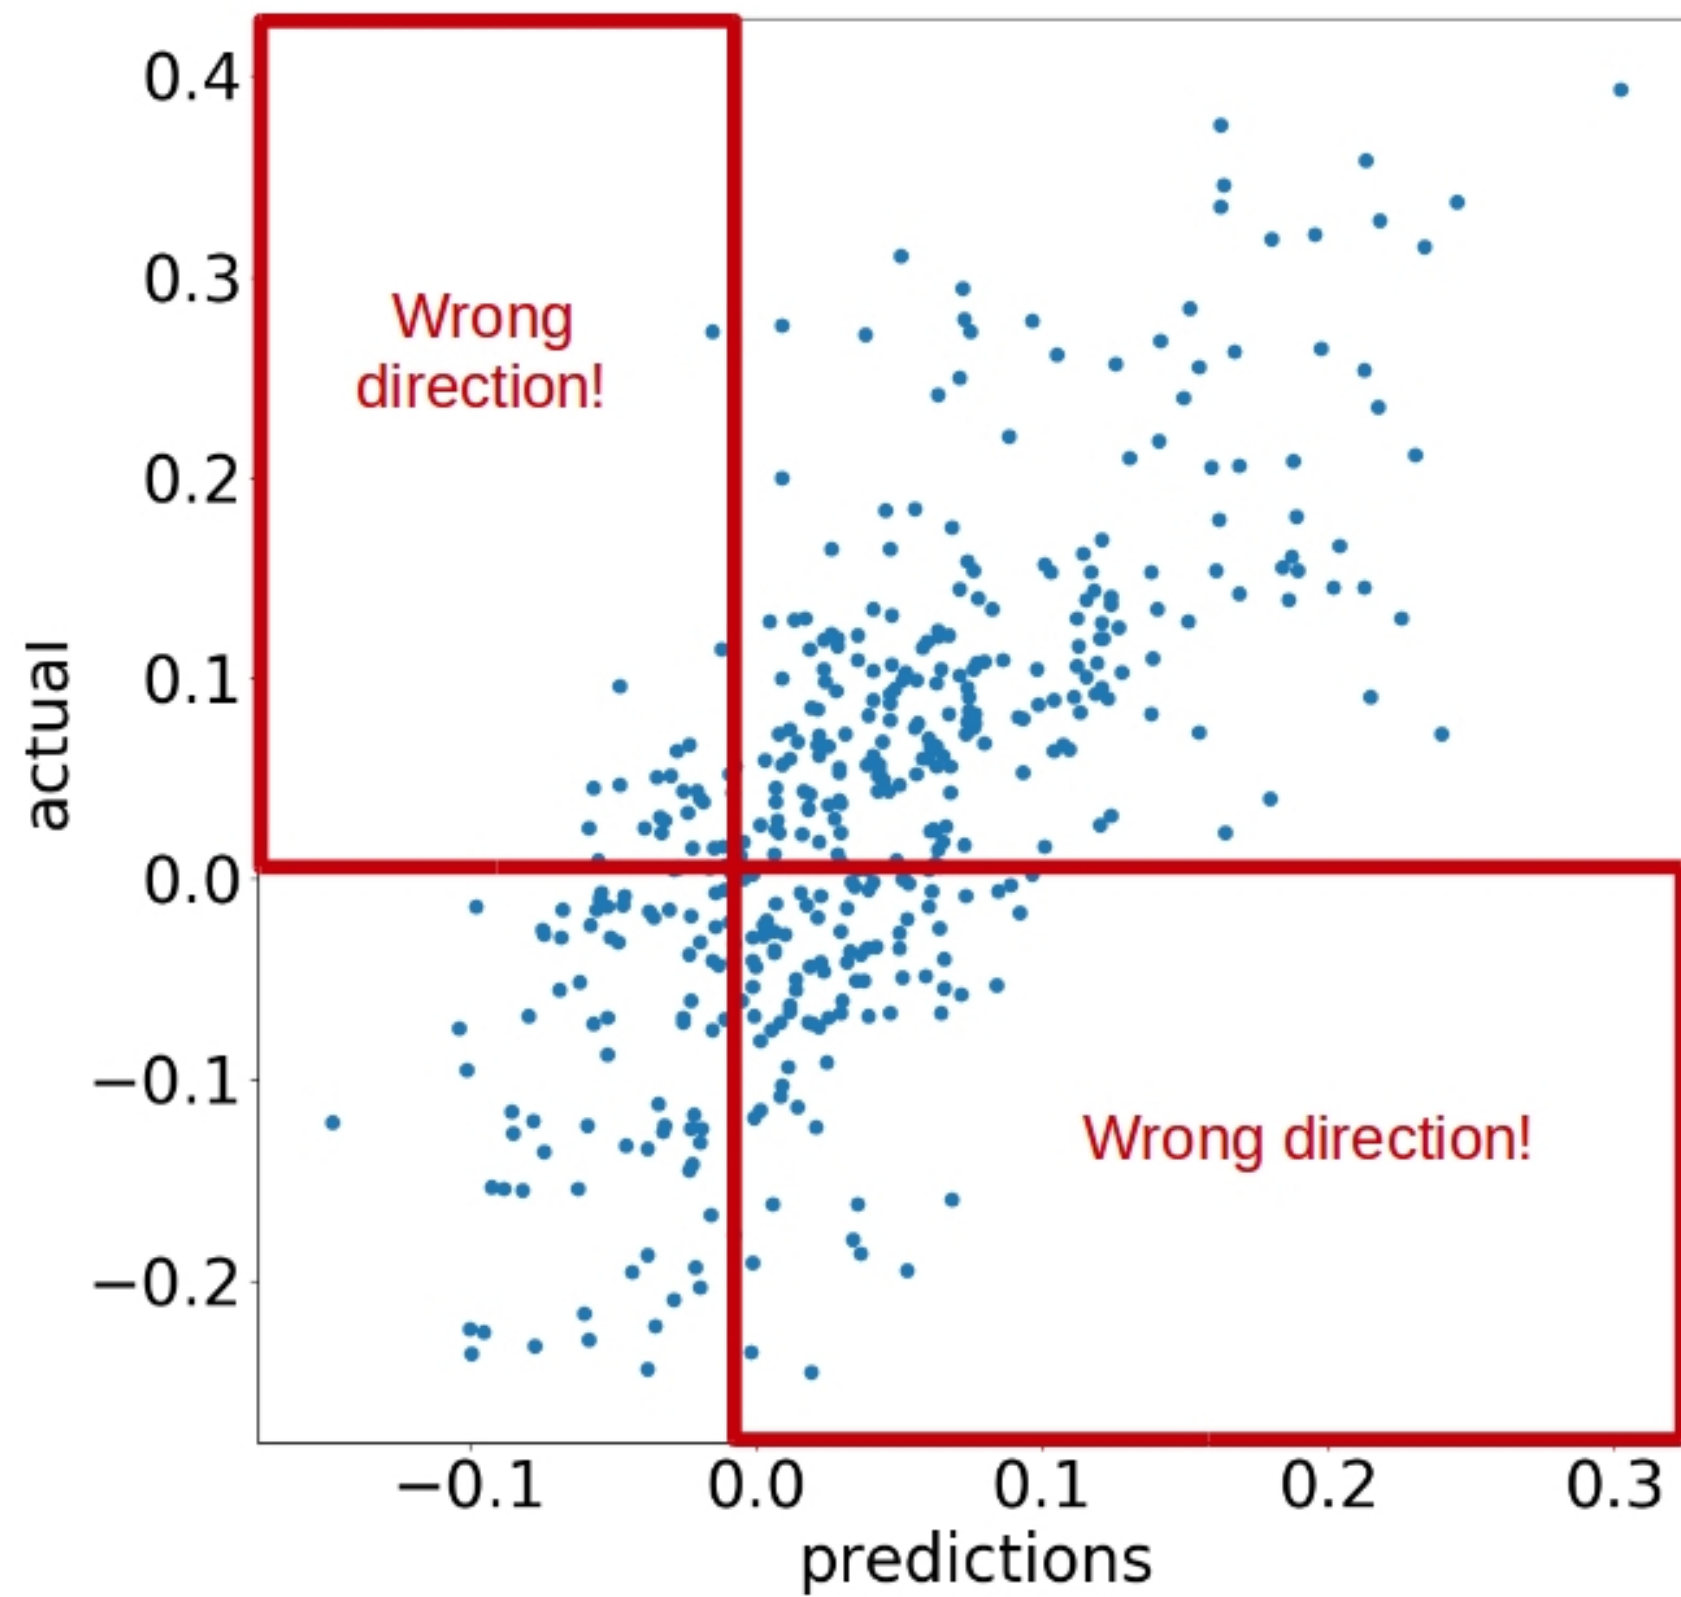


MACHINE LEARNING FOR FINANCE IN PYTHON

# Custom loss functions

Nathan George

Data Science Professor





# MSE with directional penalty

If prediction and target direction match:

- $\sum (y - \hat{y})^2$

If not:

- $\sum (y - \hat{y})^2 * \text{penalty}$



# Implementing custom loss functions

```
import tensorflow as tf
```



# Creating a function

```
import tensorflow as tf

# create loss function
def mean_squared_error(y_true, y_pred):
```



# Mean squared error loss

```
import tensorflow as tf

# create loss function
def mean_squared_error(y_true, y_pred):

    loss = tf.square(y_true - y_pred)
    return tf.reduce_mean(loss, axis=-1)
```





# Add custom loss to keras

```
import tensorflow as tf

# create loss function
def mean_squared_error(y_true, y_pred):
    loss = tf.square(y_true - y_pred)
    return tf.reduce_mean(loss, axis=-1)

# enable use of loss with keras
import keras.losses
keras.losses.mean_squared_error = mean_squared_error

# fit the model with our mse loss function
model.compile(optimizer='adam', loss=mean_squared_error)
history = model.fit(scaled_train_features, train_targets, epochs=50)
```



# Checking for correct direction

```
tf.less(y_true * y_pred, 0)
```

Correct direction:

- $\text{neg} * \text{neg} = \text{pos}$
- $\text{pos} * \text{pos} = \text{pos}$

Wrong direction:

- $\text{neg} * \text{pos} = \text{neg}$
- $\text{pos} * \text{neg} = \text{neg}$



# Using tf.where()

```
# create loss function
def sign_penalty(y_true, y_pred):
    penalty = 10.
    loss = tf.where(tf.less(y_true * y_pred, 0), \
                    penalty * tf.square(y_true - y_pred), \
                    tf.square(y_true - y_pred))
```



# Tying it together

```
# create loss function
def sign_penalty(y_true, y_pred):
    penalty = 100.
    loss = tf.where(tf.less(y_true * y_pred, 0), \
                    penalty * tf.square(y_true - y_pred), \
                    tf.square(y_true - y_pred))

    return tf.reduce_mean(loss, axis=-1)

keras.losses.sign_penalty = sign_penalty # enable use of loss with keras
```

# Using the custom loss

```
# create the model
model = Sequential()
model.add(Dense(50,
                input_dim=scaled_train_features.shape[1],
                activation='relu'))
model.add(Dense(10, activation='relu'))
model.add(Dense(1, activation='linear'))

# fit the model with our custom 'sign_penalty' loss function
model.compile(optimizer='adam', loss=sign_penalty)
history = model.fit(scaled_train_features, train_targets, epochs=50)
```



# The bow-tie shape

```
train_preds = model.predict(scaled_train_features)

# scatter the predictions vs actual
plt.scatter(train_preds, train_targets)
plt.xlabel('predictions')
plt.ylabel('actual')
plt.show()
```



MACHINE LEARNING FOR FINANCE IN PYTHON

**Create your own loss  
function!**




MACHINE LEARNING FOR FINANCE IN PYTHON

# Overfitting and ensembling

Nathan George

Data Science Professor









# Simplify your model

Complex net overfits



Simpler net prevents overfitting






# Neural network options

Options to combat overfitting:


- Decrease number of nodes
- Use L1/L2 regularisation
- Dropout
- Autoencoder architecture
- Early stopping
- Adding noise to data
- Max norm constraints
- Ensembling

# Dropout

33% dropout



no dropout





# Dropout in keras

```
from keras.layers import Dense, Dropout

model = Sequential()
model.add(Dense(500,
                input_dim=scaled_train_features.shape[1],
                activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(100, activation='relu'))
model.add(Dense(1, activation='linear'))
```



# Test set comparison

$R^2$  values on AMD without dropout:

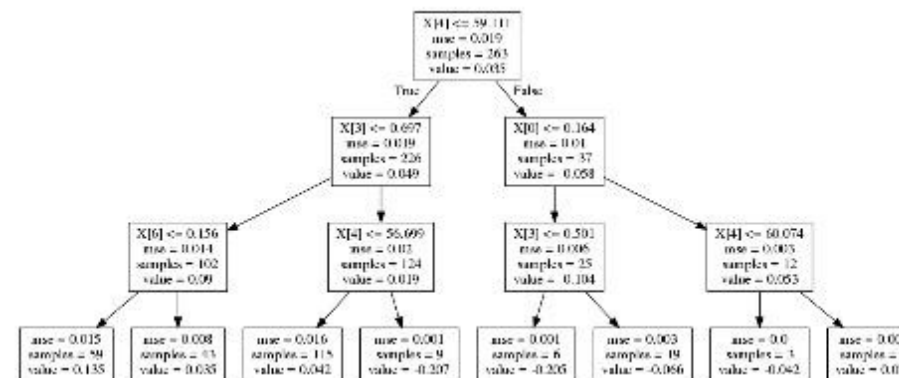
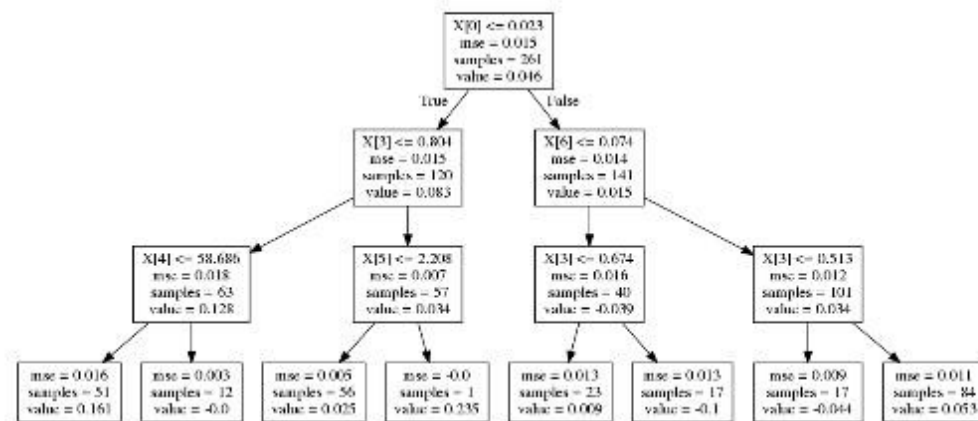


- train: 0.91
- test: -0.72

With dropout:

- train: 0.46
- test: -0.22



# Ensembling





# Implementing ensembling

```
# make predictions from 2 neural net models
test_pred1 = model_1.predict(scaled_test_features)
test_pred2 = model_2.predict(scaled_test_features)

# horizontally stack predictions and take the average across rows
test_preds = np.mean(np.hstack((test_pred1, test_pred2)), axis=1)
```





# Comparing the ensemble

Model 1  $R^2$  score on test set:

- -0.179

model 2:

- -0.148

ensemble (averaged predictions):

- -0.146



MACHINE LEARNING FOR FINANCE IN PYTHON

# Dropout and ensemble!